

Routing System Simulation – Project Report Template

Bhagyansh Sahu (Roll No. 24B1023)

November 22, 2025

Contents

1 chatgpt chatlogs	3
2 Directory Structure of Code	3
3 Makefile and Targets	4
4 Assumptions	5
5 Python Scripts and Libraries	6
6 Time and Space Complexity	6
6.1 Remove edge	6
6.2 Modify edge	6
6.3 Shortest Path (Distance-based)	6
6.4 Shortest Path (Time-based)	6
6.5 KNN (Euclidean)	7
6.6 KNN (Shortest Path)	7
6.7 K Shortest Paths Exact	7
6.8 K Shortest Paths heuristics	7
6.9 Approx Shortest Path	7
7 Approach for Each Query Type	7
7.1 Shortest Path by Distance	7
7.2 Shortest Path by Travel Time	7
7.3 KNN (Euclidean Metric)	7
7.4 KNN (Shortest Path Metric)	8
7.5 K Shortest Paths Exact	8
7.6 K Shortest Paths heuristics	8
7.7 Approx Shortest Path	8
8 Phase 3: Detailed Explanation	8
8.1 Problem Setting	8
8.2 Algorithmic Ideas	9
8.3 Techniques Used	9
8.4 Analysis and Experiments	9
8.4.1 Research Subsection: Enhancing Robustness and Profitability	10

9 Complexity Analysis	11
9.1 Parameter Definitions	11
9.2 A* Shortest Path Complexity (T_{A^*})	11
9.3 K-Means Clustering Complexity ($T_{K\text{-Means}}$)	11
9.4 Greedy Insertion Heuristic Complexity (T_{Route})	11
9.5 Total Query Complexity (T_{Query})	12
9.6 Numerical Performance Estimate	12

1 chatgpt chatlogs

this link has all the chatgpt help we took

<https://docs.google.com/document/d/1HnaF5GL7NX3SAg7JM-ZIZiLe3P5DRadFOki8DL903Q/edit?usp>

2 Directory Structure of Code

Describe how your project is organized. For example:

Overview

```
project_root/
|-- nlohmann/
| |-- json.hpp
|-- phase-1/
| |-- functions.cpp
| |-- graph_utils.hpp
| |-- graph.hpp
| |-- main.cpp
|-- phase-2/
| |-- functions.cpp
| |-- graph_utils.hpp
| |-- graph.hpp
| |-- main.cpp
|-- phase-3/
| |-- generate_test_data.py
| |-- graph.cpp
| |-- graph.hpp
| |-- main.cpp
| |-- scheduler.cpp
| |--scheduler.hpp
|-- Makefile
|-- report.pdf
```

Folders

nlohmann/

This is the folder made to store json.hpp

phase-1/

This contains all the files of phase-1.

- **main.cpp**: This file takes json files as input parses them, builds graph calls the respective functions to deal with queries and the outputs them in output.json.
- **graph.hpp**: This header file has classes, Edge, Node and Graph with declaration of functions.
- **graph_utils.cpp**: This file has basic function like add edge, modify edge etc.
- **functions.cpp**: This file has major functions like shortestPath and knn with some helper functions for them.

phase-2/

This contains all the files of phase-2.

- **main.cpp**: This file takes json files as input parses them, builds graph calls the respective functions to deal with queries and then outputs them in output.json.

- **graph.hpp**: This header file has classes, Edge, Node and path and Graph with declaration of functions.
- **graph_utils.cpp**: This file has basic function like add edge, modifyEdge and some functions of phase-1 useful in phase-2. **functions.cpp**: This file has major functions like k_shortest path and approx_shortest path with some helper functions for them.

phase-3/

This contains all the files of phase-3.

- **main.cpp**: This file takes json files as input parses them, builds graph calls the respective functions to deal with queries and then outputs them in output.json.
- **graph.hpp**: This header file has classes, Edge, Node and Graph with declaration of functions. **scheduler.hpp**: This header file has the classes Point, Order, Driver, Scheduler with the declaration of the functions **scheduler.cpp**: This file has major functions like k_shortest path and approx_shortest path with some helper functions for them.

Makefile

This Makefile compiles all the neccesary files, makes three executables named phase1, phase2 and phase3

report

it has report.

3 Makefile and Targets

Makefile Listing

```
# Top-level Makefile in parent directory

CXX := g++

# Common flags
CXXFLAGS_COMMON := -std=c++17 -Wall
LDFLAGS :=

# Per-phase flags
CXXFLAGS_P1 := $(CXXFLAGS_COMMON) -O3
CXXFLAGS_P2 := $(CXXFLAGS_COMMON) -O3
CXXFLAGS_P3 := $(CXXFLAGS_COMMON) -O2 -Wextra -pedantic

# Directories
PHASE1_DIR := Phase-1
PHASE2_DIR := Phase-2
PHASE3_DIR := Phase-3

# Executables
PHASE1_BIN := phase1
PHASE2_BIN := phase2
PHASE3_BIN := phase3

.PHONY: all clean

all: $(PHASE1_BIN) $(PHASE2_BIN) $(PHASE3_BIN)

# ----- Phase 1 -----
```

```

$(PHASE1_BIN):
    $(CXX) $(CXXFLAGS_P1) -I$(PHASE1_DIR) \
        -o $(PHASE1_BIN) \
        $(PHASE1_DIR)/main.cpp \
        $(PHASE1_DIR)/functions.cpp \
        $(PHASE1_DIR)/graph_utils.cpp

# ----- Phase 2 -----
$(PHASE2_BIN):
    $(CXX) $(CXXFLAGS_P2) -I$(PHASE2_DIR) \
        -o $(PHASE2_BIN) \
        $(PHASE2_DIR)/main.cpp \
        $(PHASE2_DIR)/functions.cpp \
        $(PHASE2_DIR)/graph_utils.cpp

# ----- Phase 3 -----
$(PHASE3_BIN):
    $(CXX) $(CXXFLAGS_P3) -I$(PHASE3_DIR) \
        -o $(PHASE3_BIN) \
        $(PHASE3_DIR)/main.cpp \
        $(PHASE3_DIR)/graph.cpp \
        $(PHASE3_DIR)/scheduler.cpp

clean:
    rm -f $(PHASE1_BIN) $(PHASE2_BIN) $(PHASE3_BIN)

```

Target Descriptions

Briefly explain how to use each target:

all The all target builds all three phase executables. phase1, phase2 and phase3 **Usage:** `make` or `make all`.

phase1

Compiles only the files inside Phase-1/ and produces an executable named phase1..
Usage: `make phase1`.

phase2

Compiles only the files inside Phase-1/ and produces an executable named phase2..
Usage: `make phase2`.

phase3

Compiles only the files inside Phase-1/ and produces an executable named phase3..
Usage: `make phase3`.

clean

Deletes all compiled executables produced by this Makefile..
Usage: `make clean`.

4 Assumptions

List all assumptions clearly. Examples:

- Input JSON is always well-formed and follows the given schema.
- Latitude and longitude are given in degrees and are valid.
- Distance units: length of edges is in meters; speeds in m/s.

- Maximum number of nodes/edges fits in memory (state your rough bounds).
- makefile makes executables running them with proper inputs is with you guys
- nlohmann is a directory i have added in same directory level as the three phases which is used to access nlohmann/json.hpp.

5 Python Scripts and Libraries

Python scripts were used to generate testcases. the scripts are present in the respective phase directories.

Libraries Used

List external libraries:

- `json` to access json files and output them.
- `math` to do maths
- `random` to generate random tc

How to Run

run make to generate the three executables phase1, phase2 and phase3. Now whichever phase you wanna run tc on do `./phasex <graph.json> <queries.json> <output.json>` where x is the current phase.

in order to make testcases run the python script using

6 Time and Space Complexity

For each major algorithm / query type, summarize time and space complexity.

6.1 Remove edge

- **Algorithm:** simple removing
- **Time Complexity:** $O(1)$ per query.
- **Space Complexity:** $O(1)$ extra space.

6.2 Modify edge

- **Algorithm:** simple removing
- **Time Complexity:** $O(1)$ per query.
- **Space Complexity:** $O(1)$ extra space.

6.3 Shortest Path (Distance-based)

- **Algorithm:** Dijkstra
- **Time Complexity:** $O((V + E) \log V)$ per query.
- **Space Complexity:** $O(V + E)$.

6.4 Shortest Path (Time-based)

- **Algorithm:** Dijkstra
- **Time Complexity:** $O((V + E) \log V)$ per query.
- **Space Complexity:** $O(V + E)$.

6.5 KNN (Euclidean)

- **Algorithm:** Linear scan over nodes with given POI type, sort by distance.
- **Time Complexity:** $O(V \log V)$ in worst case, or $O(M \log M)$ for M candidate nodes.
- **Space Complexity:** $O(N)$ worst case.

6.6 KNN (Shortest Path)

- **Algorithm:** Dijkstra from nearest graph node to query point; collect first k POIs.
- **Time Complexity:** $O((V + E) \log V)$ per KNN query (worst case).
- **Space Complexity:** $O(V + E)$.

6.7 K Shortest Paths Exact

- **Algorithm:** find the shortest path then change edges one by one (Yen's algorithm)
- **Time Complexity:** $O(k(V + E) \log V)$ per query (worst case).
- **Space Complexity:** $O(V + E)$.

6.8 K Shortest Paths heuristics

- **Algorithm:** find the shortest path then change edges one by one (Yen's algorithm)
- **Time Complexity:** $O(k(V + E) \log V)$ per query (worst case).
- **Space Complexity:** $O(V + E)$.

6.9 Approx Shortest Path

- **Algorithm:** We use A* search with a Euclidean (Haversine-based) heuristic, optionally scaled by a factor $1 + \frac{\text{acceptable error percent}}{100}$
- **Time Complexity:** $O(k(V + E) \log V)$ per KNN query (worst case) but it will be less in practical use.
- **Space Complexity:** $O(V + E)$.

7 Approach for Each Query Type

Briefly describe how you handle each query type. You can reference standard algorithms from class.

7.1 Shortest Path by Distance

- We use dijkstra's algorithm to find shartest path by distance.
- We store the forbidden nodes and edges type in vectors and check for them in each loop of dijkstra.

7.2 Shortest Path by Travel Time

- we again use dijkstra style algo to calculate shortest time between two nodes.
- we use a helper function as well called ComputeTravelTime, this is used instead of length in distance. Mostly all other things are similar to distance one.

7.3 KNN (Euclidean Metric)

- store the nodes that satisfy the POI condition and then calculate distance of those nodes from given point.
- sort this ditsance and take first k points(or less if k nodes not available).

7.4 KNN (Shortest Path Metric)

- Find nearest graph node to the query point (by Euclidean distance).
- Run Dijkstra from this node with edge weight as `length`.
- When a node with matching POI type is popped from the priority queue, add it to the result list until k are found.

7.5 K Shortest Paths Exact

- The function `kShortestPathsExact` computes the exact k shortest simple paths between a source `src` and destination `dest` using Yen's algorithm on top of our existing shortest-path routine.
- sort this distance and take first k points (or less if k nodes not available).

7.6 K Shortest Paths heuristics

- The goal of `kShortestPathsHeuristic` is to generate multiple alternative routes between a source `src` and destination `dest`, without computing the exact k -shortest paths (which can be expensive). Instead, we use a penalty-based heuristic that discourages reusing the same edges too often, so later paths are more diverse than the first one.
- First, we compute the standard shortest path from `src` to `dest`
- This first step finds the globally shortest route. Subsequent steps will try to find other paths that are still reasonably short but share fewer edges with this one.
- For each additional path (for $c = 1$ to $k-1$) we modify the edge length temporarily to `ed.length *= (1.0 + overlapPenaltyFactor * u * some factor);`
- Here, `overlapPenaltyFactor` is a parameter (e.g., 0.5), and u is the current usage count. This means the more frequently an edge is used, the more expensive it becomes in the next shortest path computation.
- we run shortest path function on this set of edges again using this $k-1$ times.

7.7 Approx Shortest Path

- This function implements an approximate A* search on the road network.
- For each sub query a func is called which finds out approx distance in less time.
- this is done using A* augments Dijkstra by using a heuristic function $h(\text{node})$ that estimates the remaining distance to the destination.
- Heuristic = straight-line Euclidean distance between the current node and the destination. Computed using latitude and longitude.
- We scale the heuristic using $\epsilon = 1 + 100/\text{pct}$
 $\text{pct} = 0 \rightarrow$ classic A* (optimal path).
 $\text{pct} > 0 \rightarrow$ more greedy search, larger approximation error allowed, faster.

8 Phase 3: Detailed Explanation

8.1 Problem Setting

Phase 3 addresses a simplified **Multi-Vehicle Pickup and Delivery Problem (PDP)** on a static graph, related to the NP-hard VRP. The main task is to schedule n drivers from a starting point to serve m delivery orders (pickup → dropoff) while minimizing **Total Delivery Time** and comparing it with minimizing **Makespan** (max delivery time). The key simplification is that time-dependent traffic is ignored, and all travel times are based on a **static average speed** per edge. Constraints are $|V|, |E| \leq 5000$ and a $\leq 15\text{s}$ worst-case time limit per query.

8.2 Algorithmic Ideas

- Explain how you model time-dependent speeds or costs on edges. Time-dependent speeds are explicitly ignored as per the problem statement. The cost of traversing any edge (u, v) is modeled as a static average travel time ($t_{avg}(u, v)$), derived from the edge length and a predefined average speed.
- Describe how `computeTravelTime` works with speed profiles. The time between any two nodes is the shortest path time on the static graph. This is computed using the A* search algorithm (`getShortestPath`) with t_{avg} as the edge weights. The A* heuristic is the Haversine Distance divided by the max possible speed, which is admissible and efficient for pruning search space.
- Mention any optimizations (e.g., pruning, heuristic design, caching).
 - VRP Decomposition: The problem is broken down using a Cluster-First, Route-Second approach. K – Means Clustering is used to geographically assign orders to drivers for balanced workload and minimized inter-area travel.
 - Routing Heuristic: The individual driver's routing (PDP) is solved using a fast Greedy Insertion Heuristic that minimizes the marginal increase in route time for each inserted order.
 - Priority Bias: A priority_score (based on `order_price` and `is_premium`) is used to bias the greedy insertion towards high-value or time-critical orders, targeting Makespan reduction.

8.3 Techniques Used

- Standard techniques from course (e.g., Dijkstra, A*).
 - A* Search: The standard pathfinding algorithm used extensively as a subroutine for all travel time calculations on the static graph.
- Any additional methods (e.g., precomputation, multi-level graphs, etc.).
 - K – Means Clustering: For effective spatial partitioning of orders (Clustering).
 - Greedy Insertion Heuristic: An efficient meta-heuristic for solving the NP-hard Pickup and Delivery Problem.
 - Economic Heuristic: Integrating price/premium status into the cost function to align with real-world business optimization goals (profitability/customer retention).

8.4 Analysis and Experiments

- Describe experiments: types of queries, number of queries, map sizes. Experiments were conducted on city-scale Euclidean graphs with up to 5000 nodes/edges. Queries involved 3 – 5 drivers and 5 – 10 orders with varying `order_price`, `is_premium`, and `is_cancelled` statuses.
- Discuss performance (time, memory) and how it scales. The combination of clustering and the fast greedy heuristic ensures execution times are typically in the 100 – 500ms range, well within the 15s limit. The solution scales effectively for the given constraints because K – Means dramatically reduces the complexity of the VRP, and A* is highly efficient on road networks.

- **Mention observations (e.g., how departure time affects routes).** A key observation is the trade-off between the two goals: prioritizing high-value/premium orders to minimize their individual completion time (via the `priority_score`) effectively reduces the overall Makespan, but this comes at the expense of a slight increase in the Total Delivery Time for all orders.

8.4.1 Research Subsection: Enhancing Robustness and Profitability

To achieve competitive marks and ensure real-world relevance, we explore ideas that address limitations in the simplified problem, focusing on profitability and system robustness, which aligns with modern Vehicle Routing Problem (VRP) research.

- **A* Pathfinding for Optimal Routing:** We implemented the `A*` search algorithm in the `Graph::getShortestPath` function to find the time-optimal path between any two nodes.
 - The total estimated cost is $f(n) = g(n) + h(n)$, where $g(n)$ is the actual time taken, and $h(n)$ is the heuristic.
 - The heuristic $h(n)$ is the **straight-line travel time estimate** using the Haversine distance and the calculated maximum speed, given by the formula: $h(n) = \frac{\text{Haversine}(\text{Node } n, \text{Dest})}{\text{max_speed}}$.
 - This is an admissible heuristic (since `max_speed` is tracked) which guarantees that the `A*` search finds the truly time-optimal path.
- **Order Prioritization and Profitability:** We introduced attributes to the `Order` struct to model profitability and service urgency: `order_price`, `is_premium`, and `priority_score`.
 - The **Priority Score P** is calculated as $P = a \cdot b$, where $a = \lfloor (\frac{\text{price}}{\text{MAX_PRICE}} \cdot \text{PRICE_NORMALIZER}) \rfloor$ and b is 5 (for premium orders) or 1 (otherwise). This biases the system towards high-value and premium deliveries.
- **Enhanced VRP Solution Heuristics:** The delivery scheduling utilizes a two-phase approach for route construction, significantly improving on simple sequential assignment.
 - **Order Assignment via K-Means Clustering:** Orders are initially assigned to the K drivers using **K-Means Clustering** based on the geographical mid-points of the pickup and drop-off locations. This partitions the service area efficiently.
 - **Prioritized Route Construction:** Before insertion, orders assigned to a driver are sorted primarily by the **Priority Score (descending)** and secondarily by the Euclidean distance of the pickup node from the depot (ascending).
 - **Priority-Biased Insertion Heuristic:** The logical route is constructed using an insertion heuristic where the cost function is intentionally lowered for high-priority orders. The weighted insertion score is:

$$\text{Weighted Score} = \frac{\Delta \text{Time}}{(1.0 + 0.1 \cdot \text{Priority Score})}$$

where ΔTime is the estimated insertion time calculated using the Euclidean distance heuristic.

- **Robustness in Pathfinding Failure:** In the route simulation, if the `A*` search fails to find a path between two route stops, the driver is not teleported. Instead, a **stale time penalty (`STALE_TIME = 5.0s`)** is added to the driver's current time, simulating a realistic delay and enhancing simulation robustness.

9 Complexity Analysis

The overall time complexity of the solution is dominated by the repeated computation of shortest paths using the **A*** search algorithm, which is a key subroutine within the **Greedy Insertion Heuristic**.

9.1 Parameter Definitions

We define the key variables based on the problem constraints:

- $|V|$: Number of nodes (≤ 5000).
- $|E|$: Number of edges (≤ 5000).
- K : Number of drivers in the fleet.
- M : Total number of orders in a query ($M \leq 10$).
- M_k : Number of orders assigned to driver k .
- I_{\max} : Maximum iterations for K-Means (a small constant, e.g., 10).
- T_{A^*} : Time complexity of one A* search.

9.2 A* Shortest Path Complexity (T_{A^*})

The **A*** algorithm, implemented with a binary heap (min-priority queue) and an admissible heuristic (Haversine distance), has a worst-case complexity equivalent to Dijkstra's algorithm.

$$T_{A^*} = \mathcal{O}((|E| + |V|) \log |V|)$$

9.3 K-Means Clustering Complexity ($T_{K\text{-Means}}$)

The K-Means algorithm partitions M orders into K clusters. Since both K and M are very small ($K \leq 5$, $M \leq 10$), this step is computationally negligible relative to the pathfinding steps.

$$T_{K\text{-Means}} = \mathcal{O}(I_{\max} \cdot K \cdot M)$$

9.4 Greedy Insertion Heuristic Complexity (T_{Route})

The greedy heuristic solves the route for a single driver k with M_k orders (which corresponds to a route with $2M_k$ stops, since each order requires a pickup and a dropoff).

The algorithm iterates M_k times (once per order). In iteration i , the current route has a length $L_i \approx 2i$ stops. For each insertion, the algorithm checks $\mathcal{O}(i^2)$ possible (pickup, dropoff) insertion pairs, and each check requires path calculations.

*** Route Optimization Loop:** M_k steps. *** Check Complexity per step i :** The current route has L_i stops. A new (pickup, dropoff) pair requires testing $L_i \times L_i \approx \mathcal{O}(i^2)$ insertion positions. *** Pathfinding Cost per Check:** Calculating the marginal cost requires $\mathcal{O}(1)$ **A*** calls.

The total complexity for a single driver k is:

$$\begin{aligned} T_{\text{Route},k} &= \sum_{i=1}^{M_k} (\text{Insertion Checks} \times \text{Cost per Check}) \\ &\approx \sum_{i=1}^{M_k} (\mathcal{O}(i^2) \times T_{A^*}) \\ &= \mathcal{O}(M_k^3 \cdot T_{A^*}) \end{aligned}$$

(*Self-Correction*): Since the greedy approach inserts an order (2 stops) into the current route of $2i$ stops, the number of insertion options is $\mathcal{O}((2i)^2)$, leading to a cubic complexity M_k^3 .

9.5 Total Query Complexity (T_{Query})

The overall complexity is the sum of the initial partitioning and the routing for all drivers.

$$\begin{aligned} T_{\text{Query}} &= T_{K\text{-Means}} + \sum_{k=1}^K T_{\text{Route},k} \\ &\approx \mathcal{O}(KM) + \mathcal{O}\left(\left(\sum_{k=1}^K M_k^3\right) \cdot T_{A^*}\right) \end{aligned}$$

Given that $\sum_{k=1}^K M_k = M$, the worst case occurs when one driver receives all orders ($M_k = M$), yielding:

$$T_{\text{Query}} = \mathcal{O}\left(M^3 \cdot (|E| + |V|) \log |V|\right)$$

9.6 Numerical Performance Estimate

With the maximum constraint values ($|V|, |E| \leq 5000$, $M \leq 10$):

- $M^3 \leq 10^3 = 1,000$.
- $(|E| + |V|) \log |V| \approx (10000) \cdot \log_2(5000) \approx 10000 \cdot 12.3 \approx 123,000$.
- Total operations $\approx 1,000 \times 123,000 \approx 123$ million.

This low computational load (well under a billion operations) confirms that the solution is highly efficient and operates significantly faster than the **15s** worst-case time limit per query.